KV-weights are all you need for skipless transformers

Nils Graef[*]
OpenMachine

Abstract

He and Hofmann [1] detailed a skipless transformer without the V and P (post-attention projection) linear layers, which reduces the total number of weights. However, this scheme is only applicable to MHA (multi-head attention) [2], but not for MQA (multi-query attention) [3] and GQA (grouped-query attention) [4]. The latter schemes are used by many popular LLMs such as Llama 2, Mistral, Mixtral, PaLM, and Gemma [5, 6, 7, 8, 9]. Therefore, this micro-paper [10] proposes mathematically equivalent versions that are suitable for MQA and GQA. For example, removing Q and P from a skipless version of Mistral-7B would remove 15% of its weights, and thus reduce its compute and memory complexity. Watch our explainer video [11] and see [12, 13] for code and more transformer tricks.

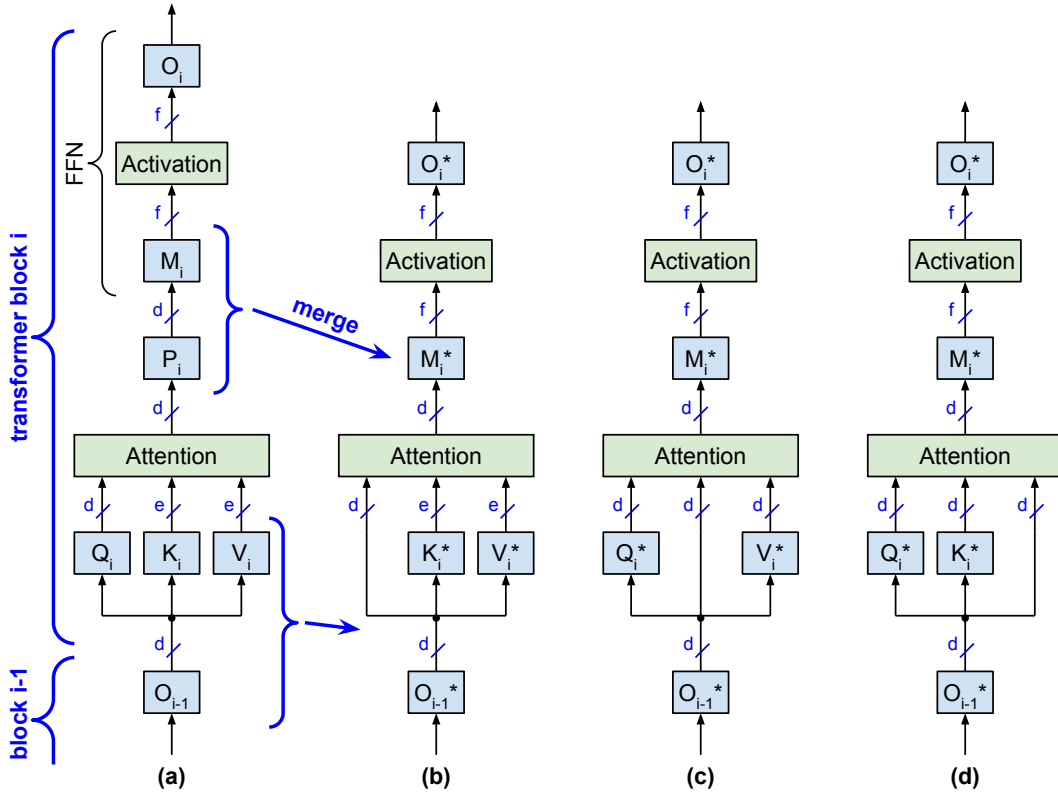## 1 Vanilla transformer without skip connections

Figure 1: (a) Skipless vanilla transformer; equivalent versions with (b) Q and P merged into the FFN (feedforward network); (c) K and P merged into FFN; (d) V and P merged into FFN. $\mathbf{M}_i^*, \mathbf{Q}_i^*, \mathbf{K}_i^*, \mathbf{V}_i^*, \mathbf{O}_{i-1}^*$ are defined in table 1.

---

[*]info@openmachine.ai

He et al. [14] have shown how transformers without skip connections and normalization (Figure 1(a)) can be trained successfully. Removing skip connections and normalization allows us to merge linear layers in a mathematically identical way as shown in Figures 1(b) to (d). This reduces the number of weights without changing the functionality as follows:

- Figure 1(b) is mathematically identical to Figure 1(a) and eliminates $2d^2$ weights per transformer block by merging $\mathbf{P}_i$ into $\mathbf{M}_i^*$ and $\mathbf{Q}_i$ into $\mathbf{O}_{i-1}^*$.

- For MHA where $e = d$, Figures 1(c) and (d) are mathematically identical to Figure 1(a) and eliminate $2d^2$ weights per transformer block by merging $\mathbf{P}_i$ into $\mathbf{M}_i^*$ and $\mathbf{K}_i$ or $\mathbf{V}_i$ into $\mathbf{O}_{i-1}^*$.

- This requires that $\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i$ are invertible (i.e. nonsingular). It is extremely rare that a square matrix with random values is not invertible [15] (which requires its determinant to be exactly 0).

Figure 1 uses the following dimensions and weight matrices, based on the type of attention:

- d: embedding dimension

- e: $e = d$ for MHA. For MQA, $e = d/n_{heads}$. And for GQA, $e = d \cdot n_{kv\_heads}/n_{heads}$.

- f: hidden dimension of the FFN. $f = 4d$ in the vanilla transformer; Shazeer [3] uses $f > 4d$. For models that use a GLU variant [16] (such as Llama and Mistral), the effective f' for the first linear layer M is $f' = 2f$, because the GLU variant uses two linear layers that are combined (via pointwise multiplication) with a non-linear activation function.

- $\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i, \mathbf{P}_i$: The weight matrices of the linear layers for query, keys, values, and the post-attention projection of transformer block i.

- $\mathbf{M}_i, \mathbf{O}_i$: The weight matrices of the FFN input and output linear layers.
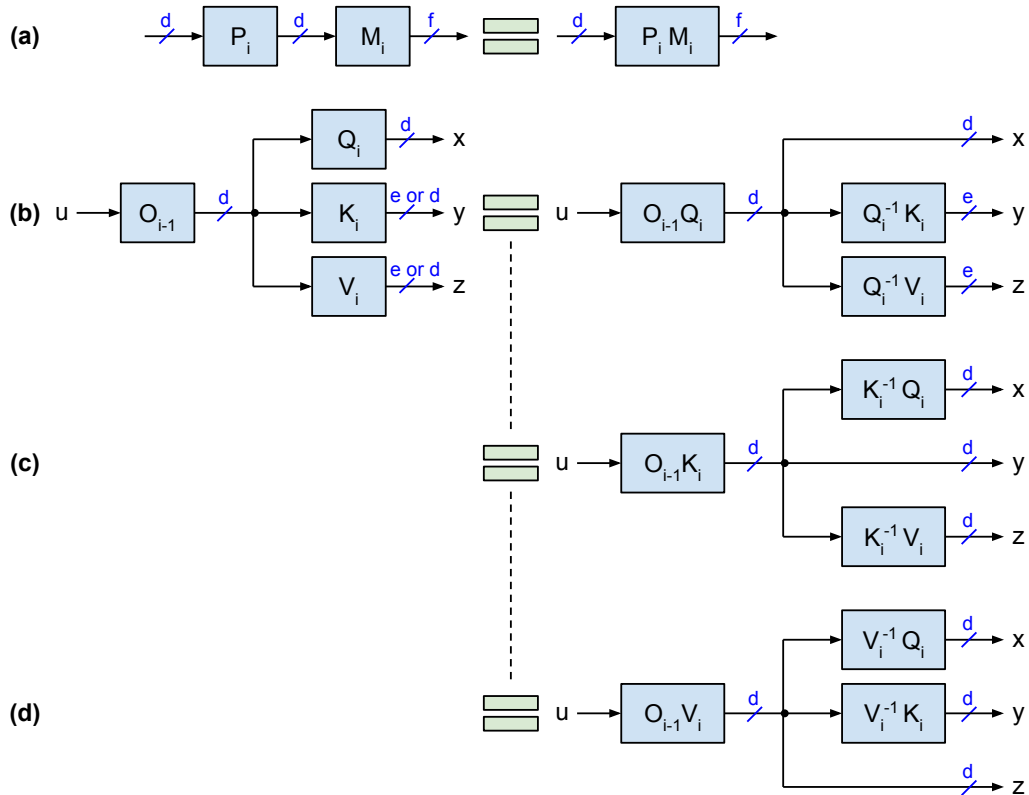

Figure 2: (a) Merging P and M; (b) eliminating Q; (c) eliminating K; (d) eliminating V.

Figure 2 details how the linear layers are merged:

- Figure 2(a) shows how the two linear layers with weight matrices $\mathbf{P}_i$ and $\mathbf{M}_i$ are collapsed and replaced by a single linear layer with weight matrix $\mathbf{M}_i^* = \mathbf{P}_i\mathbf{M}_i$, which eliminates d^2 weights.

- Figure 2(b) illustrates how to merge $\mathbf{Q}_i$ into the preceding $\mathbf{O}_{i-1}$-matrix, which eliminates d^2 weights and requires $\mathbf{Q}_i$ to be invertible. Note that $\vec{y} = \vec{u}\mathbf{O}_{i-1}(\mathbf{Q}_i\mathbf{Q}_i^{-1})\mathbf{K}_i = \vec{u}\mathbf{O}_{i-1}\mathbf{K}_i$ and $\vec{z} = \vec{u}\mathbf{O}_{i-1}(\mathbf{Q}_i\mathbf{Q}_i^{-1})\mathbf{V}_i = \vec{u}\mathbf{O}_{i-1}\mathbf{V}_i$.
- For MHA where $e = d$, $\mathbf{K}_i$ can be removed as shown in Figure 2(c), which eliminates d^2 weights. Note that $\vec{x} = \vec{u}\mathbf{O}_{i-1}(\mathbf{K}_i\mathbf{K}_i^{-1})\mathbf{Q}_i = \vec{u}\mathbf{O}_{i-1}\mathbf{Q}_i$ and $\vec{z} = \vec{u}\mathbf{O}_{i-1}(\mathbf{K}_i\mathbf{K}_i^{-1})\mathbf{V}_i = \vec{u}\mathbf{O}_{i-1}\mathbf{V}_i$. This requires that $\mathbf{K}_i$ is invertible.
- For MHA where $e = d$, $\mathbf{V}_i$ can be removed as shown in Figure 2(d), which eliminates d^2 weights. Note that $\vec{x} = \vec{u}\mathbf{O}_{i-1}(\mathbf{V}_i\mathbf{V}_i^{-1})\mathbf{Q}_i = \vec{u}\mathbf{O}_{i-1}\mathbf{Q}_i$ and $\vec{y} = \vec{u}\mathbf{O}_{i-1}(\mathbf{V}_i\mathbf{V}_i^{-1})\mathbf{K}_i = \vec{u}\mathbf{O}_{i-1}\mathbf{K}_i$. This requires that $\mathbf{V}_i$ is invertible.

Table 1 specifies how the new weight matrices ($\mathbf{M}_i^*, \mathbf{Q}_i^*, \mathbf{K}_i^*, \mathbf{V}_i^*, \mathbf{O}_{i-1}^*$) of Figure 1 are calculated from the original ones. For the first transformer block ($i = 1$), we use the input embedding instead of $\mathbf{O}_{i-1}$ (because there is no $\mathbf{O}_{i-1}$ for $i = 1$).

	Figure 1(b)	Figure 1(c)	Figure 1(d)
$\mathbf{O}_{i-1}^*$	$\mathbf{O}_{i-1}\mathbf{Q}_i$	$\mathbf{O}_{i-1}\mathbf{K}_i$	$\mathbf{O}_{i-1}\mathbf{V}_i$
$\mathbf{Q}_i^*$	1 (eliminated)	$\mathbf{K}_i^{-1}\mathbf{Q}_i$	$\mathbf{V}_i^{-1}\mathbf{Q}_i$
$\mathbf{K}_i^*$	$\mathbf{Q}_i^{-1}\mathbf{K}_i$	1 (eliminated)	$\mathbf{V}_i^{-1}\mathbf{K}_i$
$\mathbf{V}_i^*$	$\mathbf{Q}_i^{-1}\mathbf{V}_i$	$\mathbf{K}_i^{-1}\mathbf{V}_i$	1 (eliminated)
$\mathbf{M}_i^*$	$\mathbf{P}_i\mathbf{M}_i$	$\mathbf{P}_i\mathbf{M}_i$	$\mathbf{P}_i\mathbf{M}_i$

Table 1: How to calculate the new weight matrices from the original ones for Figure 1.

2 Parallel transformer without skip connections

Similar to the parallel transformer [17], Figure 3 shows parallel versions of Figures 1(b) to (d). Here, "parallel" refers to having the attention (including its linear layers) in parallel to the FFN.
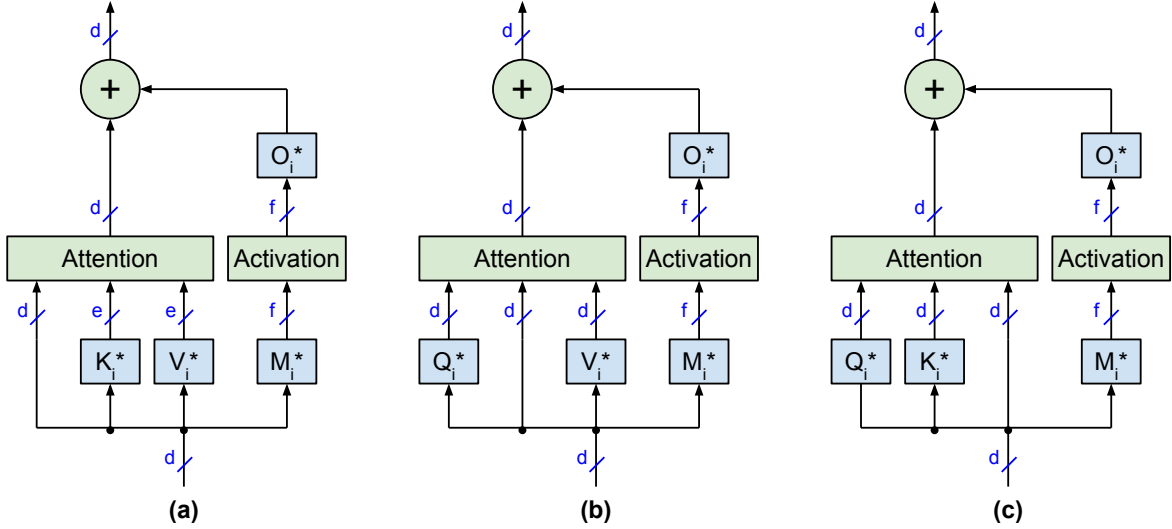

Figure 3: Parallel skipless transformers (a) without Q and P; (b) without K and P; (c) without V and P.

Figures 3(b) and (c) require that $e = d$, so they are only suitable for MHA, but not for MQA and GQA. Figure 3(a) is suitable for MHA, MQA, and GQA. Figure 3(c) is identical to the simplified transformer proposed in [1].

3 Related work

In addition to [1], our work is related to the lossless weight compression for back-to-back linear layers presented in [18], and FlashNorm's weight fusion [19].

4 Examples

The table below lists the configurations and weight counts for Pythia-6.9B and Mistral-7B. For a skipless version of Mistral-7B we would save 15% of weights after merging the Q and P linear layers into the FFN layers. For a batch 1 system that is limited by memory bandwidth, these 15% weight savings can speed up inference by 1.17x during the autoregressive next-token-generation phase, see the table below.

Parameter	Pythia-6.9B	Mistral-7B	Notes
Parallel attention/FFN?	parallel	serial	[17]
MHA, MQA, or GQA?	MHA	GQA	[2, 3, 4]
dim (aka d)		4,096	embedding dimension
n_layers		32	number of layers
n_heads		32	number of heads
n_kv_heads	32	8	number of KV-heads
e (output dim. of K, V)	4,096	1,024	e = d * n_kv_heads / n_heads
FFN type	MLP	MLP with SwiGLU	[16]
FFN hidden_dim	16,384	14,336	FFN hidden dimension
vocab_size	50,400	32,000	vocabulary size
Number of weights (calculated from above parameters):			
Q+P weights per layer		33,554,432	2 * dim * dim
K+V weights per layer	33,554,432	8,388,608	2 * dim * dim / n_heads * n_kv_heads
FFN weights per layer	134,217,728	176,160,768	(2 or 3) * dim * hidden_dim
Input+output embed.	412,876,800	262,144,000	2 * dim * vocab_size
Total weights:	6.9B	7.2B	
Weight savings and speedup after removing Q and P:			
Total w/o Q+P weights:	5.8B	6.2B	total after removing Q and P
Weight savings:	**16%**	**15%**	
Possible speedup:	**1.19x**	**1.17x**	assumes batch size 1

5 Experiments

Refer to [12] for Python code that demonstrates the numerical equivalency of the weight reduction illustrated in Figures 1(b) and 2(b). The code also confirms that all square matrices of Mistral-7B are invertible.

6 Conclusion

A novel approach to optimizing skipless transformers by eliminating the query (Q) and post-attention projection (P) linear layers is presented. This mathematical equivalent weight fusion offers savings in computational cost, memory, and energy consumption by reducing the number of weights.

Recently published skipless transformers such as [20, 14] could be retrofitted post-training with the lossless weight fusion presented here. Skipless transformers with normalization layers could first eliminate the normalization layers by fine-tuning as described in [21, 22] and then apply the weight fusion described in our work.

Our work extends a recent trend in neural network design toward architectural parsimony, in which unnecessary components are removed to create more efficient models. Notable examples include RMSNorm's simplification of LayerNorm by removing mean centering [23], PaLM's elimination of bias parameters [8], decoder-only transformers' omission of the encoder stack [24], FlashNorm's elimination of normalization weights [19], and Slim Attention's elimination of the V-cache [25]. This trend is rooted in the revolutionary shift initiated by the original Transformer model, which replaced traditional recurrence and convolutions with a more streamlined architecture [2].

Because skipless transformers are not very popular right now, future work should investigate whether removing P and Q (or K or V) is also beneficial for transformers with normalization and skip connections as illustrated in Figure 4. Adding normalization and skip connections again could simplify and speed up training relative to skipless transformers.
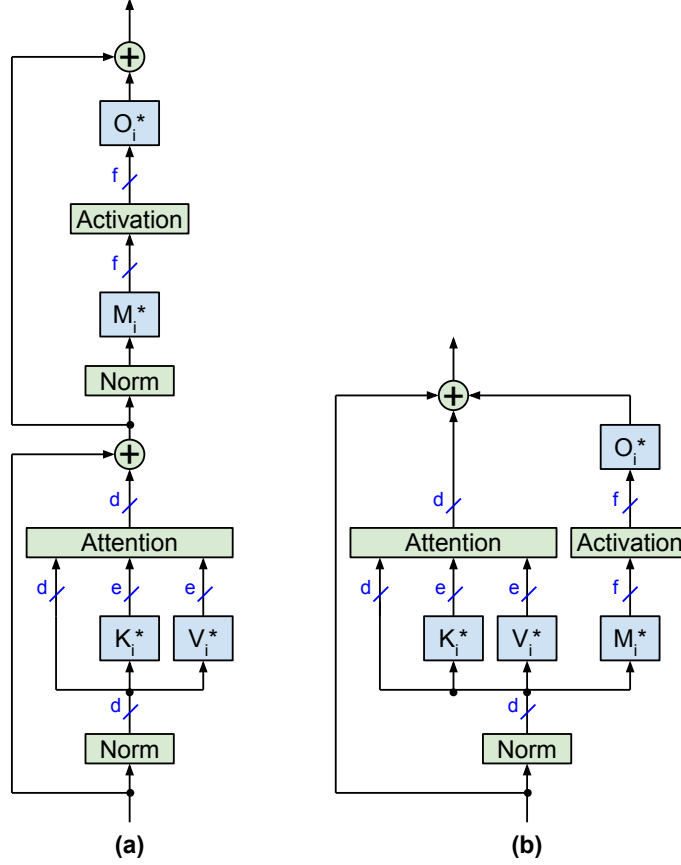
Figure 4: (a) Transformer block without Q and P; (b) version with parallel attention / FFN.

Acknowledgments

References

[1] Bobby He and Thomas Hofmann. Simplifying Transformer Blocks. November 2023. *arXiv:2311.01906*.

[2] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. June 2017. *arXiv:1706.03762*.

[3] Noam Shazeer. Fast Transformer Decoding: One Write-Head is All You Need. November 2019. *arXiv:1911.02150*.

[4] Joshua Ainslie, James Lee-Thorp, Michiel de Jong, Yury Zemlyanskiy, Federico Lebrón, and Sumit Sanghai. GQA: Training generalized multi-query transformer models from multi-head checkpoints. May 2023. *arXiv:2305.13245*.

[5] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, Dan Bikel, Lukas Blecher, Cristian Canton Ferrer, Moya Chen, Guillem Cucurull, David Esiobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, Anthony Hartshorn, et al. Llama 2: Open foundation and fine-tuned chat models. July 2023. *arXiv:2307.09288*.

[6] Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, Lélio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El Sayed. Mistral 7B. October 2023. *arXiv:2310.06825*.

[7] Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, Gianna Lengyel, Guillaume Bour, Guillaume Lample, Lélio Renard Lavaud, Lucile Saulnier, Marie-Anne Lachaux, Pierre Stock, Sandeep Subramanian, Sophia Yang, Szymon Antoniak, Teven Le Scao, Théophile Gervet, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El Sayed. Mixtral of Experts. January 2024. *arXiv:2401.04088*.

[8] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, Parker Schuh, Kensen Shi, Sasha Tsvyashchenko, Joshua Maynez, Abhishek Rao, Parker Barnes, Yi Tay, Noam Shazeer, et al. PaLM: Scaling language modeling with Pathways. April 2022. *arXiv:2204.02311*.

[9] Gemma Team, Google DeepMind. Gemma: Open Models Based on Gemini Research and Technology. 2024.

[10] Frank Elavsky. The Micro-Paper: Towards cheaper, citable research ideas and conversations. February 2023. *arXiv:2302.12854*.

[11] Christopher Marquand with audio generated by Notebook LM. Explainer video for paper 'Removing weights for skipless transformers'. Oct 2025.

[12] OpenMachine. Transformer tricks. 2024. URL https://github.com/OpenMachine-ai/transformer-tricks.

[13] Nils Graef. Transformer tricks: Precomputing the first layer. February 2024. *arXiv:2402.13388*.

[14] Bobby He, James Martens, Guodong Zhang, Aleksandar Botev, Andrew Brock, Samuel L Smith, and Yee Whye Teh. Deep transformers without shortcuts: Modifying self-attention for faithful signal propagation. February 2023. *arXiv:2302.10322*. And ICLR 2023.

[15] Wikipedia. Invertible matrix, 2024. Accessed Mar-2024.

[16] Noam Shazeer. GLU Variants Improve Transformer. February 2020. *arXiv:2002.05202*.

[17] Ben Wang and Aran Komatsuzaki. GPT-J-6B: A 6 billion parameter autoregressive language model. 2021. *Github repo*.

[18] Nils Graef. MatShrink: Lossless weight compression for back-to-back linear layers. 2025. URL https://github.com/OpenMachine-ai/transformer-tricks/blob/main/doc/matShrink.pdf.

[19] Nils Graef, Andrew Wasielewski, and Matthew Clapp. FlashNorm: fast normalization for LLMs. July 2024. *arXiv:2407.09577*.

[20] Yiping Ji, James Martens, Jianqiao Zheng, Ziqin Zhou, Peyman Moghadam, Xinyu Zhang, Hemanth Saratchandran, and Simon Lucey. Cutting the skip: Training residual-free transformers. September 2025. *arXiv:2510.00345*.

[21] Stefan Heimersheim. You can remove GPT2's LayerNorm by fine-tuning. September 2024. *arXiv:2409.13710*.

[22] Luca Baroni, Galvin Khara, Joachim Schaeffer, Marat Subkhankulov, and Stefan Heimersheim. Transformers don't need LayerNorm at inference time: Scaling LayerNorm removal to GPT-2 XL and the implications for mechanistic interpretability. July 2025. *arXiv:2507.02559*.

[23] Biao Zhang and Rico Sennrich. Root mean square layer normalization. October 2019. *arXiv:1910.07467*.

[24] Peter J Liu, Mohammad Saleh, Etienne Pot, Ben Goodrich, Ryan Sepassi, Lukasz Kaiser, and Noam Shazeer. Generating Wikipedia by summarizing long sequences. January 2018. *arXiv:1801.10198*.

[25] Nils Graef and Andrew Wasielewski. Slim attention: cut your context memory in half without loss – K-cache is all you need for MHA. March 2025. *arXiv:2503.05840*.